

Quaternion Neural Networks for 3D Sound Source Localization in Reverberant Environments: Implementation using First Order Ambisonics

Roberto Aureli, ID 1757131
Riccardo Caprari, ID 1743168
Gianmarco Fioretti, ID 1762135

Neural Networks, Winter 2020

Contents

1	Introduction	2
2	Dataset	3
3	Architecture	4
4	Our work	5
4.1	Project architecture	5
5	Results	8
6	Conclusions	10

1 Introduction

This project addresses the **3D Sound Event Localization and Detection Task** in reverberant environments with a quaternion neural network (deep neural network with quaternion input features extracted from the acoustic intensity vector).

As mentioned above, the proposed model performs both sound localization and sound event detection and subsequent classification. In particular it follows the architecture described in 2020 paper (TO ADD REFERENCE) which allows to estimate the three-dimensional direction of arrivals (DOA), in addition, it has been modified in order to be capable of detecting sound events and estimating the corresponding sources (eleven different classes provided by the development dataset, check for Table 1).

Many recent works have proven that deep quaternion neural networks are able to improve localization performances dramatically, especially in reverberant and noisy environments, thanks to spatial harmonic decomposition which permits to exploit the intrinsic correlation of the ambisonics signal components. One of the main aspect to be considered is the input features to be passed to the network (TO FILL).

Our main work consisted in adapting the provided code to the new dataset, fine-tuning the model's hyper-parameters, applying the three metrics defined in the 2019 paper (TO ADD REFERENCE) (**SED**, **DOA**, **SELD**), implementing the direction of arrival estimation in spherical coordinates, and computing the confidence intervals on the model's final errors as defined in the 2018 paper (TO ADD REFERENCE) in order to add also a statistical evaluation on final results.

2 Dataset

The network is trained with the **TAU Spatial Sound Events 2019** dataset which provides four-channel directional microphone recordings of stationary point sources.

It is a balanced dataset, it indeed consists of eleven classes, each with twenty examples that were randomly split into five sets with an equal number of examples for each class, in addition, it was divided into four cross-validation split.

The maximum number of simultaneously overlapping sources are two. Moreover, in order to improve the performance over new sound events, and to make a more realistic scenario, natural ambient noise collected in the recording locations was added to the synthesized recordings in the dataset such that the average SNR of the sound events was 30 dB.

Table 1: TAU Spatial Sound Events Classes

Sound class	Index
knock	0
drawer	1
clearthroat	2
phone	3
keysDrop	4
speech	5
keyboard	6
pageturn	7
cough	8
doorslam	9
laughter	10

3 Architecture

The network proposed in the 2020 paper (TO ADD REFERENCE) involves a series of convolutional layers in the quaternion domains, they are composed of several filter kernels which allow learning inter-channel features, with subsequent activation functions (TO FILL) and max-pooling functions. The output of this series are properly reshaped and fed to bidirectional quaternion recurrent layers. This first part of the architecture is depicted in Figure 1




Figure 1: First part of the Network Architecture

As already pointed out on Chapter 1 the network has been modified accordingly, the final part of architecture therefore becomes as in Figure 2, so that network is able to perform a multi-classification for the SED task (one for each class: 1 to indicate detection, 0 otherwise) and multi-regressions for DOA task (one for the azimuth angle, one for the elevation angle and one for the distance).

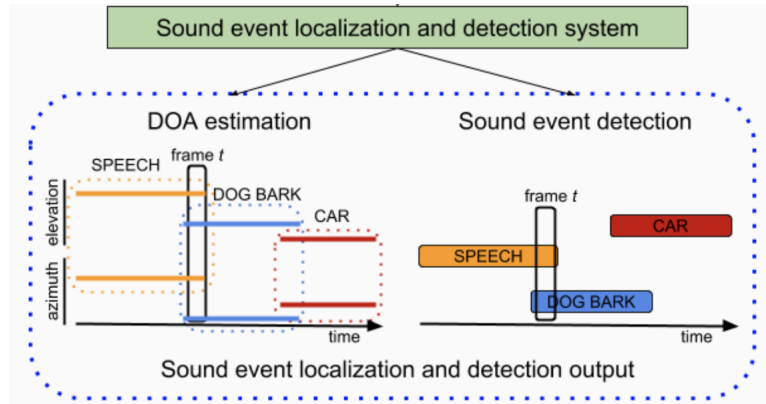


Figure 2: Last part of the Network Architecture

4 Our work

The objective of our work is to build a working SELDnet based network that works with First Order Ambisonics data sets. In particular, we are going to extend SELDnet, adding capabilities to both support pre-existing data sets (ansim, resim, etc.) and the FOA one in a smart, modular, performing way. Other metrics have been added like the SELD score, mainly used in the paper (rif) outcomes evaluation and a tiny library for a graphical representation of the results. Now we'll proceed with a more technical analysis of the work, listing all the features added.

4.1 Project architecture

We decided to maintain a parameters manager file to configure the SELDnet, like its input dataset with specific values, training, evaluation and test process combinations and many other settings which can be found in *parameter.py*. We can divide our project in 3 main classical phases: extraction, training and test.

Extraction

Before starting with the training, we have to extract features from the new FOA data set. This delicate process makes the difference between a successful training in terms of evaluation metrics. For this reason we decided to introduce a new *FeatureClass*, which will differ from the old one in some parameters and procedures. In fact, as we experienced, introducing a new similar class results in a way more readable, modifiable and modular approach. We noticed, in fact, that different data sets can be similar from a structure point of view and totally different with respect to their parameters, like record settings. For this reason a new full-dedicated extractor results in a wise and fast choice.

As we can observe in scheme 3, *batch_feature_extraction.py* has the task of managing the extraction process, differing function calls with respect to data set type. The input is divided into two folders, *foa_dev* and *meta-data_dev*, which contains audio records stored with a *.wav* format and their corresponding labels.

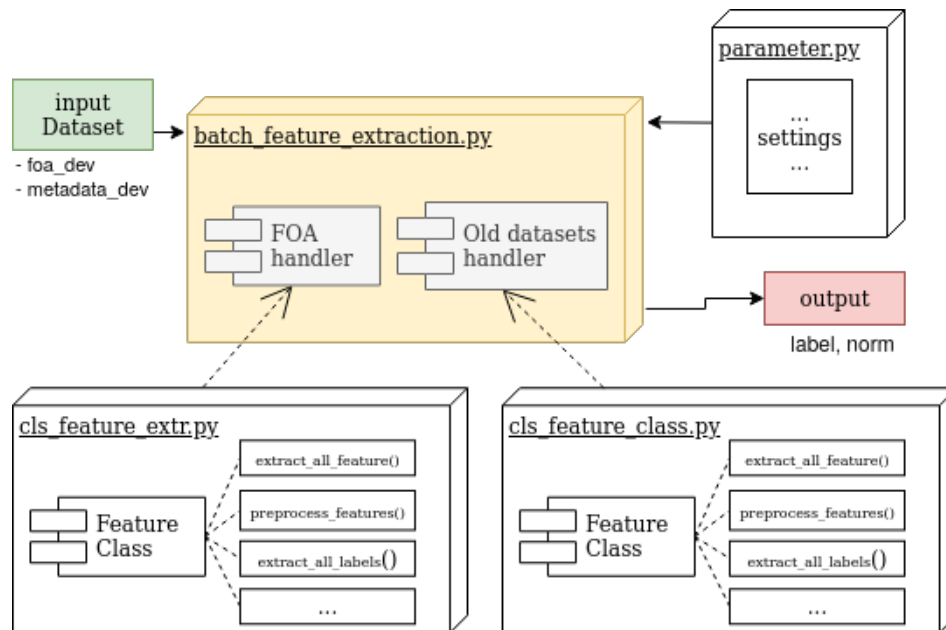


Figure 3: Extraction process.

Let's now zoom in the new *cls_feature Extr.py* extractor and highlight the main differences with the old *FeatureClass*.

Training

Test

Summary

5 Results

Speak about metrics and comment results...




Figure 4: Loss function over 100 epochs.




Figure 5: Sed, Doa and Seld scores over 100 epochs.




Figure 6: Sed error rate and f1-score over 100 epochs.

6 Conclusions

TO ADD CONCLUSIONS

TO ADD REFERENCES